

Cleaved hemagglutinin in ESM3:

Can the model treat HA1 and HA2 as one molecule
joined by a disulfide bridge?

my_extendedSampling / softmax / bridge_test

2026-10-09

Abstract

`zero_polymer` is an influenza hemagglutinin precursor (HA0). In the mature protein it is cut once, after Arg329 in mature numbering, into HA1 and HA2, which stay together through an interchain disulfide between Cys14 of HA1 and Cys137 of HA2. The question was whether ESM3 can be given that two-chain molecule and handle it. The mechanical answer is yes: the stock tokenizer accepts the chain-break character, but no benefit of declaring the break shows up consistently across decoding step counts. The structural answer is no, with caveats: no assembled construct folds well enough to pass the geometry gate, and the interchain Cys14–Cys137 distance is 96 Å (cleaved) and 130 Å (no break), against a bonding criterion of 4.5 Å. Two known-disulfide controls then showed that the distance criterion itself is weaker than I first said: it has some sensitivity but no specificity. This report records the results, the corrections made along the way, the limitations, and a prioritised to-do list.

Contents

1	Summary	1
2	What was built	2
3	Method	2
4	Results	3
4.1	The controls	3
4.2	HA constructs	3
4.3	Step sweep	3
4.4	Findings	4
4.5	Disulfide calibration	5
5	Corrections made during the session	5
6	Limitations	6
6.1	Of the measurements	6
6.2	Of the scope	6
6.3	Of the session itself	6
7	To do	7
7.1	Recommended next step	7
7.2	Prioritised list	7
7.3	Housekeeping	8
8	Reproducing the runs	8

1 Summary

- **Mechanics work.** `model.encode(ESMProtein(sequence=...))` accepts 1; 551 positions go in and 551 come out; `to_protein_complex()` recovers two chains. No change to the model code was needed.
- **No consistent benefit of the chain-break.** At 256 steps HA1’s mean pLDDT is 51.4 fused to HA2, 63.2 with the break and 65.2 alone, which looked like a benefit. Across 64, 128 and 256 steps the break-minus-fused difference is -5.2 , $+2.4$, $+11.8$: the sign is not stable, and the earlier claims that the break helps HA1 and its Cys14 anchor are withdrawn.
- **Step count is a large, non-monotone effect.** HA1 alone scores 48.4, 67.0, 65.2 and 59.8 at 64, 128, 256 and 329 steps. Each step count is a deterministic but effectively independent draw, and draws differ by more than most effects claimed so far. Only sign-stable differences are findings.
- **HA2 alone folds well; every assembly degrades it.** HA2 alone passes the gate at all three step counts (pLDDT 88–94, no clashes) and is 17–29 points more confident than HA2 in any assembly. No assembled construct, and no HA1-alone run, passes the gate at any step count.
- **The interchain bridge is not formed** in either assembled construct, by a margin (at least $10\times$ the criterion at every step count) that no threshold subtlety can change.
- **The 4.5 Å criterion is not a reliable bond detector.** BPTI: 3/3 known bonds inside it. Crambin: 1/3 inside it, and the closest pair in the molecule is a non-native one at 3.33 Å.
- **Not tested:** the trimer, any experimental reference structure, an interchain interface measure, partner identity for HA.

2 What was built

All of it is on branch `claude/esm3-disulfide-polymer-chains-rwvpoq`.

Commit	Content
46bccd8	<code>tests/references.py</code> : chains, bridge, numbering helpers
1f35fbf, fcb86aa, caeb436	<code>tests/test_cleaved_complex.py</code> , three revisions
4056d91, 346b3de	BPTI and crambin controls; known-disulfide measurement
086e2cd, 1b0d3ce, 6b34313,	<code>bridge_test/DEVLOG_bridge_test.txt</code>
8efcd2c	
87dc571	<code>bridge_test/DEVLOG_calibration.txt</code>

The construct. Chain one is mature HA1 (329 aa, ends `...QTR`); chain two is HA2 (221 aa, starts `GIFGAI`); the 16-residue signal peptide is dropped, so 550 residues in total. Both chains are slices of `ZERO_POLYMER`, never retyped, and `_check()` asserts the junction, the lengths, and that both anchors are cysteines.

A numbering trap worth keeping. The HA landmarks are in *mature* numbering, which runs 16 behind the index into `ZERO_POLYMER`. Read naively, `ZERO_POLYMER[328]` (“Arg329”) is T and `ZERO_POLYMER[13]` (“Cys14”) is V. Neither raises. `mature_to_index()` and `chain_index()` exist so the conversion is never done by hand.

3 Method

Route into the model. The sampler’s usual path (`ExtendedProtein` with `predict_attention` on soft `sequence_probs`) cannot carry a chain break: `expand()` maps through a 25-letter vocabulary with no `|`, and the encoder zero-pads over token 31. The stock `encode/predict_protein` route runs the real tokenizer and has none of these problems. Two shortcuts that look right are wrong: `chain_id=` is consumed only by geometric attention, which the soft path masks and zeroes, so it is a silent no-op; and `sequence_id=` is an attention separator, so distinct ids would make the chains ignore each other.

What the decoder returns. Backbone plus an inferred $C\beta$, and no side-chain sulfur at any residue. So SG-SG is unavailable and $C\beta$ - $C\beta$ distance is the only disulfide proxy. $C\beta$ is not predicted: `infer_cbeta()` derives it from N/CA/C and writes NaN at every glycine. The chain-break position is excluded from all $C\beta$ statistics.

The instrument.

- *Geometry gate:* $pTM \geq 0.40$, and no $C\beta$ - $C\beta$ pair closer than $3.0 \text{ \AA}$ at sequence separation ≥ 2 .
- *Bond criterion:* $C\beta$ - $C\beta \leq 4.5 \text{ \AA}$.
- *Controls:* protein G (fold control, no cysteines); BPTI and crambin (three known disulfides each).

4 Results

4.1 The controls

Reference	steps	pTM	pLDDT	% ≥ 70	min $C\beta$ ($\text{\AA}$)	clashes
protein G (56 aa)	56	0.7596	93.24	100.0	3.68	0
BPTI (58 aa)	58	0.7545	92.89	98.3	3.35	0
crambin (46 aa)	46	0.5816	83.31	89.1	3.33	0

All three pass the gate, so the harness folds, and the choice of $3.0 \text{ \AA}$ as the clash line is borne out empirically (confident folds have their closest pair at 3.3 – $3.7 \text{ \AA}$).

4.2 HA constructs

	cleaved	mature	HA1 alone	HA1 alone	HA2 alone
positions / steps	551 / 256	550 / 256	329 / 256	329 / 329	221 / 221*
pTM	0.2322	0.2614	0.4642	0.4133	0.5674
pLDDT, all residues	63.89	59.20	65.16	59.83	93.93
% residues ≥ 70	45.5	31.8	37.1	18.5	97.7
pLDDT, HA1 (% ≥ 70)	63.22 (38.6)	51.42 (11.2)	65.16 (37.1)	59.83 (18.5)	—
pLDDT, HA2 (% ≥ 70)	64.90 (55.7)	70.78 (62.4)	—	—	93.93 (97.7)
min $C\beta$ - $C\beta$ ($\text{\AA}$)	1.50	1.17	2.01	0.44	4.14
clashes	15	9	7	16	0
gate	fail	fail	fail	fail	pass
Cys14-Cys137 $C\beta$ ($\text{\AA}$)	95.67	129.74	—	—	—
pLDDT at Cys14 / Cys137	73.6 / 61.3	22.5 / 80.9	—	—	—

The first run, at one decoding step, gave pTM 0.1827 and a bridge distance of 38.55 Å with a 1.64 Å closest cysteine pair. It is not tabulated because it measured a structure the model did not believe in; see Section 5. *HA2 alone was requested at 256 steps and clamped to 221 by the script (a chain cannot have more steps than positions), so it has been seen at one step count only.

Reproducibility. Cleaved and mature were each run twice, with different script versions, on the same machine, and gave identical pTM (0.2322 and 0.2614) and identical bridge distances (95.67 and 129.74 Å). HA2 alone was run twice as well (the second time by the clamp above) with identical output to every printed digit. Decoding at temperature zero looks deterministic here. This is three repeat pairs on one machine, not a proof.

4.3 Step sweep

	64 steps	128 steps	256 steps	
pLDDT, HA1 in cleaved	56.16	64.32	63.22	
pLDDT, HA1 fused	61.35	61.90	51.42	
pLDDT, HA1 alone	48.43	67.01	65.16	
pLDDT, HA2 in cleaved	68.49	63.27	64.90	
pLDDT, HA2 fused	68.94	71.40	70.78	
pLDDT, HA2 alone*	93.67	88.47	93.93	
pLDDT at Cys14, cleaved	31.7	58.0	73.6	
pLDDT at Cys14, fused	41.1	69.9	22.5	
clashes, cleaved / fused	22 / 9	19 / 2	15 / 9	
clashes, HA1 alone / HA2 alone	17 / 0	5 / 0	7 / 0	
bridge C β -C β , cleaved (Å)	98.8	137.9	95.7	
bridge C β -C β , fused (Å)	164.3	48.0	129.7	
Contrast at equal steps	64	128	256	sign
HA1 pLDDT: break – fused	–5.2	+2.4	+11.8	mixed
HA1 pLDDT: alone – cleaved	–7.7	+2.7	+1.9	mixed
Cys14 pLDDT: break – fused	–9.4	–11.9	+51.1	mixed
HA2 pLDDT: alone – cleaved	+25.2	+25.2	+29.0	+ at all
HA2 pLDDT: alone – fused	+24.7	+17.1	+23.2	+ at all
HA2 pLDDT: break – fused	–0.5	–8.1	–5.9	– at all

*HA2 alone is capped at its 221 residues, so its 256-step column is the 221-step run. Computed by `bridge_test/sweep_contrasts.py`; $n = 3$, so means and spreads are descriptive and no p-value is claimed.

4.4 Findings

The chain-break’s benefit to HA1 is a 256-step effect. Break minus fused is –5.2, +2.4, +11.8 at 64, 128, 256 steps. The +11.8 comes from fused HA1 dipping to 51.42 at 256 steps (61.35 and 61.90 at the other two), while cleaved HA1 stays at 56–64. The Cys14 anchor result is the same artefact: at 64 and 128 steps the break gives the *lower* anchor confidence. Neither claim passes the sign-at-every-step rule.

The clash topology differs. At 256 steps mature has fewer clashes (9 against 15) but they sit at $|i - j|$ of 80–159, so distant segments are superimposed; cleaved’s are local ($|i - j| \leq 16$). Across the sweep this holds at two of three step counts (fused is long-range at 128 and 256 but not at 64; cleaved has one long-range clash, at 182, at 128 steps), so it is suggestive, not clean.

HA2 is far more confident alone than in any assembly. HA2 alone minus HA2 in the cleaved complex is +25.2, +25.2, +29.0, and against the fused one +24.7, +17.1, +23.2: the same sign at every step count, with spreads of 2–4 points on differences of 17–29. This is the one robust contrast in the study. The break costs HA2 a little (−0.5, −8.1, −5.9, negative throughout, but of the same size as the step noise).

HA1 alone against HA1 in the complex is undetermined: −7.7, +2.7, +1.9. At 64 steps HA1 alone is unfolded (0.0% of residues above 70). Whether the complex is right in a way pLDDT cannot see, for example in the interchain placement, is untested; the assembled pTM is too low to resolve it.

Step count is a large, non-monotone effect. HA1 alone scores 48.43, 67.01, 65.16 and 59.83 at 64, 128, 256 and 329 steps; it rises, plateaus and falls. Fused HA1 drops 10 points between 128 and 256 steps while cleaved HA1 does not. Which cysteines pair also moves with steps alone (one:64–76 goes from 3.92 Å to 26.64 Å between 329 and 256 steps).

HA1 is the hard half in every construct. Alone it fails the gate at both step counts tried: 16 clashes at 329 steps (minimum $C\beta$ separation 0.44 Å) and 7 at 256 steps (2.01 Å); it fails at 64 and 128 too.

The bridge is not formed, at any step count. Cleaved gives 98.8, 137.9, 95.7 Å and fused 164.3, 48.0, 129.7 Å, against a 4.5 Å criterion. No assembled run passes the gate either (cleaved 22, 19, 15 clashes; fused 9, 2, 9, the 2-clash run failing on pTM), and HA1 alone fails at every step count tried.

4.5 Disulfide calibration

Protein	Bond	$C\beta$ – $C\beta$ (Å)	rank	within 4.5 Å
BPTI	Cys14–Cys38	3.60	1/15	yes
BPTI	Cys5–Cys55	3.66	2/15	yes
BPTI	Cys30–Cys51	4.45	3/15	yes
crambin	Cys4–Cys32	3.79	2/15	yes
crambin	Cys3–Cys40	7.75	8/15	no
crambin	Cys16–Cys26	10.42	12/15	no
crambin	Cys16–Cys40 (<i>non-native</i>)	3.33	1/15	yes
BPTI	Cys51–Cys55 (<i>non-native, closest</i>)	6.15	4/15	no

BPTI separates cleanly: the last true bond is at 4.45 Å and the first false pair at 6.15 Å. Crambin fails twice over. It misses two real bonds, and it puts a non-native pair inside the criterion. Partner recovery (known pairs among the three closest) is 3/3 for BPTI and 1/3 for crambin. A distance threshold is therefore necessary for a bond and not sufficient.

5 Corrections made during the session

These are recorded because several first readings were wrong.

1. **One decoding step.** The first verdict (“bridge not formed, 38.55 Å”) came from a script that used the default `num_steps=1`, which samples all 551 structure tokens at once. It scored pTM 0.18 with a 1.64 Å “contact” shorter than a C–C bond. The verdict was worthless. The gate now runs before any verdict.

2. **A control that did not control.** Protein G validated the fold but has no cysteines, so it could not validate the measurement. BPTI and crambin were added for that reason.
3. **Two overstated claims, retracted.** The calibration script printed “4.5 Å is sound; HA2’s disulfides are not formed” from BPTI alone. That tested sensitivity on one protein and specificity not at all. Separately, HA1-alone pairs at 3.51 and 3.92 Å had been described as “real disulfide geometry”; by the crambin result they could be non-native.
4. **A pre-registered prediction failed.** I expected the trimer signature (both chains confident, assembly low). The result was asymmetric instead: HA2 adequate, HA1 poor.
5. **“HA2’s presence helps HA1” retracted.** It rested on HA1 alone at 329 steps against HA1 in the complex at 256. At matched steps HA1 alone is 65.16 against 63.22. This report had listed that confound as its first limitation, and the rerun resolved it against the claim. I had also said the premise “ESM3 should not treat the chains as separate” was confirmed; what is confirmed is narrower (an unbroken fusion is wrong, and the break repairs it).
6. **Three more findings withdrawn after the step sweep.** The chain-break benefit to HA1, its benefit to the Cys14 anchor, and the narrower “break repairs fusion” that replaced the first claim all rested on 256 steps. I stated each with more confidence than one step count could support. The sweep also exposed a bug in my own contrast script, which kept an old pre-pLDDT result file for the 256-step runs; an empty column gave it away.
7. **Small bugs.** The chain-break position was admitted into the $C\beta$ statistics; labels were wrong for the break-free reference; a results file was overwritten by a generic filename.

6 Limitations

6.1 Of the measurements

- **No ground truth.** There is no experimental structure in the loop, so no RMSD or TM-score. pLDDT and pTM are the model’s own confidence; crambin’s pLDDT of 83 coexists with wrong partners.
- **$C\beta$ – $C\beta$ is a proxy.** The decoder emits no sulfur, and $C\beta$ is inferred from the backbone. A bond criterion built on it inherits any backbone error, and glycines (which have no $C\beta$) are invisible to the clash count.
- **Three step counts per construct.** Enough to show that the sign of several contrasts flips, not enough to characterise the step dependence. HA2 alone is capped at 221 steps. It is unknown whether steps or tokens decoded per step is the operative variable.
- **One run per condition.** The three repeat pairs agree exactly, but no condition has been run across seeds or checkpoints, and nothing here gives error bars.
- **pTM depends on length,** so comparing it between a 46-residue and a 550-residue chain is weak. Only same-construct comparisons are clean.
- **Calibration is two proteins.** BPTI and crambin are small and well-ordered. Neither says what the criterion does in a 550-residue fold.
- **Sequences were supplied, not independently verified.** The BPTI and crambin sequences and pairings came with sources. I checked them for internal consistency (length, cysteine positions, one-to-one pairing, crambin’s Pro22/Leu25 isoform) but not against UniProt directly.

- **Only esm3-open.** No other checkpoint was tried.

6.2 Of the scope

- **Monomer only.** HA is a trimer. A single HA1+HA2 protomer may have no self-contained fold, and a low score may then be correct. Not tested.
- **No verified native disulfide list for HA,** so partner recovery cannot be run on HA at all. Positive bridge claims about HA are not possible yet.
- **The energy machinery cannot take a chain break.** The sampler’s soft-probability path has no way to carry token 31, so none of the U_{am} work has been applied to the cleaved form.
- **Glycans, the fusion peptide, and everything outside the sequence** are absent by construction.

6.3 Of the session itself

- For most of the session I could not run code: the shell was blocked until auto mode was switched off in this final turn. Every script was therefore written unexecuted and validated by your runs. One consequence: `test_cleaved_complex.py` still carries an unused `present_atoms()` helper that I reintroduced by mistake.
- One commit message (87dc571) has a stray line that refers to non-existent session notes. The file content is fine; the message cannot be amended without rewriting history.

7 To do

7.1 Recommended next step

Put the measurements on a footing of experimental ground truth, using the short proteins whose structures are known, and then ask whether the bridge can be held or imposed. The scripts exist (`bridge_test/`) and are tested on synthetic data only; nothing has been fetched or run yet. In order:

1. **Fetch experimental entries** for protein G, BPTI and crambin with `fetch_pdb.py`. An entry is kept only if its own sequence matches the reference, so no entry identifier is trusted. The cloud session’s network policy blocks the PDB, so run this where it is reachable.
2. **Compare predictions with experiment** with `compare_to_pdb.py`: CA RMSD, an approximate TM-score, contact precision and recall at 8 Å, and for each known disulfide the experimental SG–SG and $C\beta$ – $C\beta$ against the predicted $C\beta$ – $C\beta$. This also tests the supplied pairings (is each pair really bonded in the file?) and shows what the experiment says about crambin’s non-native 3.33 Å pair.
3. **Round trip** (`roundtrip_bridge.py`): encode the experimental coordinates to structure tokens and decode them. If the known bonds survive, the representation can hold them and the failure is one of prediction; if they loosen, no prompt can help.
4. **Coordinate prompt**: give only the backbone of the two cysteines (then with neighbours, then every known bond) and generate. ESM3 has no residue-level disulfide input; coordinates are the only channel, and they do reach the geometric attention. Decision rules were fixed in advance in the devlog.

HA is attempted only if imposition works on the short proteins, and its geometry would then have to come from an experimental HA entry. If imposition through ESM3 fails, the fallbacks are a restrained refinement outside ESM3, or pointing the sampler’s structure-token energy at tokens taken from a bonded structure.

7.2 Prioritised list

1. **Experimental comparison, round trip and coordinate prompt** (above).
2. **Per-residue pLDDT at the known-bond residues** in the BPTI and crambin runs. If crambin’s Cys3, Cys16, Cys26 and Cys40 are low-confidence, the misses are uncertainty and a pLDDT filter may repair the criterion; if confident, the placement is wrong. The calibration block does not print this yet.
3. **Report specificity in the calibration verdict:** the count of in-range non-native pairs next to the count of in-range known bonds.
4. **Interface contacts, cleaved against fused**, once the contact metric is validated against experiment on the short proteins.
5. **A verified native disulfide list for HA**, from a named structure entry, asserted at import the way the controls are. Required before any partner-recovery claim about HA.
6. **The trimer:** one|two|one|two|one|two, 1653 positions. The biologically correct question, and expensive.
7. **Condition on a real HA structure** and ask whether ESM3 *maintains* the bridge rather than discovers it. A different and more tractable question.
8. **Only if the cleaved form is needed in the energy machinery:** inject a one-hot at token 31 in the padded region, set the structure token at the break to `STRUCTURE_CHAINBREAK_TOKEN` in `_default()`, and exclude the break site from `mutate()`.

7.3 Housekeeping

- At session end, append the two devlogs to the main one:


```
cat bridge_test/DEVLOG_bridge_test.txt \
    bridge_test/DEVLOG_calibration.txt >> DEVLOG.txt
```
- Give results files run-identifying names; the generic default overwrote the protein G control once.
- Remove the dead `present_atoms()` helper.
- `bridge_test/summarize_runs.py` tabulates every result file from its content; use it instead of copying numbers by hand.

8 Reproducing the runs

From softmax/code:

```
python tests/test_cleaved_complex.py --ref protein_g --steps 56
python tests/test_cleaved_complex.py --ref bpti      --steps 58
python tests/test_cleaved_complex.py --ref crambin   --steps 46
```

```
python tests/test_cleaved_complex.py --ref cleaved --steps 256
python tests/test_cleaved_complex.py --ref mature --steps 256
python tests/test_cleaved_complex.py --ref one --steps 329
python tests/test_cleaved_complex.py --ref one --steps 256
python tests/test_cleaved_complex.py --ref two --steps 221
```

Output files are in softmax/code/bridge_test/.